

Lab Tutorial 2

Course Project Part 2

Ahmed Elbagoury
ahmed.elbagoury@uwaterloo.ca

University of Waterloo

February 10, 2016

Concepts

Inter-Process Communication

Interrupt Handling

Timing Services

Interprocess Communication (IPC)


Requirements:

- ▶ Message-based, asynchronous
- ▶ Messages are carried in “envelopes”

Procedure:

- ▶ Process allocates memory for envelope
- ▶ Process writes data into envelope
- ▶ Process invokes message API
`send_message(proc_id, envelope)`
- ▶ Other process invokes message API
`msg_t *env = receive_message()`
and blocks if no message is available

`msg_t` (an envelope)



IPC: Process Message Queues

Where do we refer to messages that haven't been read yet?

IPC: Process Message Queues

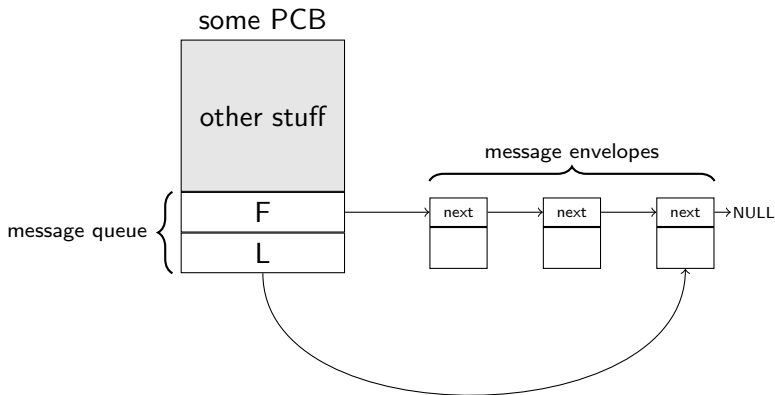
Where do we refer to messages that haven't been read yet?

We put them in a queue, tied to the PCB.

IPC: Process Message Queues

Where do we refer to messages that haven't been read yet?

We put them in a queue, tied to the PCB.



IPC: send_message()

```
void send_message(uint32 receiving_pid, msg_t *env) {  
    atomic(on);  
    set sender_procid, destination_procid;  
    pcb_t *receiving_proc = get_pcb_from_pid(receiving_pid);  
    enqueue env onto the msg_queue of receiving_proc;  
    if (receiving_proc->state is BLOCKED_ON_RECEIVE) {  
        set receiving_proc state to ready;  
        rpq_enqueue(receiving_proc);  
    }  
    atomic(off);  
}
```

Do you think send_message would ever block?

IPC: send_message()

```
void send_message(uint32 receiving_pid, msg_t *env) {
    atomic(on);
    set sender_procid, destination_procid;
    pcb_t *receiving_proc = get_pcb_from_pid(receiving_pid);
    enqueue env onto the msg_queue of receiving_proc;
    if (receiving_proc->state is BLOCKED_ON_RECEIVE) {
        set receiving_proc state to ready;
        rpq_enqueue(receiving_proc);
    }
    atomic(off);
}
```

Do you think send_message would ever block? **Nope, never!**


IPC: receive_message()

```
msg_t * receive_message() {  
    atomic(on);  
    while (current_process msg_queue is empty) {  
        set current_process state to BLOCKED_ON_RECEIVE;  
        release_processor();  
    }  
    msg_t *env = dequeue current_process msg queue;  
    atomic(off);  
    return env;  
}
```

Note: this version blocks, you will need a non-blocking one as well

Interrupt Handling

- ▶ Interrupts are hardware messages that need immediate action
- ▶ Interrupts invoke pre-registered procedures that “interrupt” currently executing code.



Interrupt Handler: Design

Requirements:

- ▶ Interrupts must be handled by processes called i-processes
- ▶ i-processes must be OS processes
(i.e., they can receive messages, use APIs)

Design:

- ▶ i-processes are scheduled by the interrupt handler
 - ▶ are always ready to run, but never in the ready queue
- ▶ i-processes can never block when invoking a kernel primitive
 - ▶ Primitives which can block need non-blocking alternatives
- ▶ Each i-process has a PCB just like other processes

Interrupt Handler: Example

Here is some pseudo C-code for a generic interrupt handler routine.

```
__asm void Timer0_IRQHandler {  
    save the context of the current_process;  
    switch the current_process with timer_i_process;  
    load the timer_i_process context;  
    call the timer_i_process C function;  
    invoke the scheduler to pick next to run process;  
    restore the context of the newly picked process;  
}
```

Interrupt Handler: Example


Here is some pseudo C-code for a generic interrupt handler routine.

```
--asm void Timer0_IRQHandler {  
    save the context of the current_process;  
    switch the current_process with timer_i_process;  
    load the timer_i_process context;  
    call the timer_i_process C function;  
    invoke the scheduler to pick next to run process;  
    restore the context of the newly picked process;  
}
```

Note: in-case you haven't noticed, this will be written in assembly

Timing Service: Design

- ▶ Timing services are fundamental to RTX
- ▶ The `timer_i_process` receives messages to deliver after a delay (i.e., `delayed_send(PID, env, delay)` API), which is non-blocking!
- ▶ After the delay expires, the i-process forwards the message
- ▶ The `timer_i_process` maintains requests in sorted list:



Timing Service: I-Process

At each clock tick (i.e., interrupt), the `timer_i_process`:

- ▶ Increments `current_time`
- ▶ Calls `receive_message` repeatedly to retrieve new requests (non-blocking)
- ▶ If there are new requests, it adds them to the queue (maintaining sorted order)
- ▶ Checks if any timing requests have expired
 - ▶ If yes, send the message to the destination

Timing Service: I-Process Example

Here is some pseudo code for the timer_i_process.

```
void timer_i_process () {  
    // get pending requests  
    while (pending messages to i-process) {  
        insert envelope into the timeout queue;  
    }  
    while (first message in queue timeout expired) {  
        msg_t *env = dequeue(timeout_queue);  
        int target_pid = env->destination_pid;  
        //forward msg to destination  
        send_message(target_pid, env);  
    }  
}
```


Questions

Do you have any questions?